

# index

March 7, 2022

## 1 Logistic Regression in scikit-learn

### 1.1 Introduction

Generally, the process for fitting a logistic regression model using scikit-learn is very similar to that which you previously saw for `statsmodels`. One important exception is that scikit-learn will not display statistical measures such as the p-values associated with the various features. This is a shortcoming of scikit-learn, although scikit-learn has other useful tools for tuning models which we will investigate in future lessons.

The other main process of model building and evaluation which we didn't to discuss previously is performing a train-test split. As we saw in linear regression, model validation is an essential part of model building as it helps determine how our model will generalize to future unseen cases. After all, the point of any model is to provide future predictions where we don't already know the answer but have other informative data ( $\mathbf{X}$ ).

With that, let's take a look at implementing logistic regression in scikit-learn using dummy variables and a proper train-test split.

### 1.2 Objectives

You will be able to:

- Fit a logistic regression model using scikit-learn

### 1.3 Import the data

```
[1]: import pandas as pd

df = pd.read_csv('titanic.csv')
df.head()
```

```
[1]: PassengerId  Survived  Pclass  \
0             1         0        3
1             2         1        1
2             3         1        3
3             4         1        1
4             5         0        3
```

```
Name      Sex   Age  SibSp  \
```

0		Braund, Mr. Owen Harris	male	22.0	1
1	Cummings, Mrs. John Bradley (Florence Briggs Th...		female	38.0	1
2		Heikkinen, Miss. Laina	female	26.0	0
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)		female	35.0	1
4		Allen, Mr. William Henry	male	35.0	0

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

#### 1.4 Define X and y

Note that we first have to create our dummy variables, and then we can use these to define X and y.

```
[2]: df = pd.get_dummies(df, drop_first=True)
print(df.columns)
df.head()
```

```
Index(['PassengerId', 'Survived', 'Pclass', 'Age', 'SibSp', 'Parch', 'Fare',
      'Name_Abbott, Mr. Rossmore Edward',
      'Name_Abbott, Mrs. Stanton (Rosa Hunt)', 'Name_Abelson, Mr. Samuel',
      ...,
      'Cabin_F G63', 'Cabin_F G73', 'Cabin_F2', 'Cabin_F33', 'Cabin_F38',
      'Cabin_F4', 'Cabin_G6', 'Cabin_T', 'Embarked_Q', 'Embarked_S'],
      dtype='object', length=1726)
```

```
[2]: PassengerId  Survived  Pclass   Age  SibSp  Parch    Fare \
0             1         0       3  22.0     1     0   7.2500
1             2         1       1  38.0     1     0  71.2833
2             3         1       3  26.0     0     0   7.9250
3             4         1       1  35.0     1     0  53.1000
4             5         0       3  35.0     0     0   8.0500

      Name_Abbott, Mr. Rossmore Edward  Name_Abbott, Mrs. Stanton (Rosa Hunt) \
0                                     0                                     0
1                                     0                                     0
2                                     0                                     0
3                                     0                                     0
4                                     0                                     0

      Name_Abelson, Mr. Samuel  ...  Cabin_F G63  Cabin_F G73  Cabin_F2 \
0                               0  ...          0           0          0
1                               0  ...          0           0          0
```

2	0	...	0	0	0	0
3	0	...	0	0	0	0
4	0	...	0	0	0	0

  

	Cabin_F33	Cabin_F38	Cabin_F4	Cabin_G6	Cabin_T	Embarked_Q	Embarked_S
0	0	0	0	0	0	0	1
1	0	0	0	0	0	0	0
2	0	0	0	0	0	0	1
3	0	0	0	0	0	0	1
4	0	0	0	0	0	0	1

[5 rows x 1726 columns]

Wow! That's a lot of columns! (Way more than is useful in practice: we now have columns for each of the passengers names. This is an example of what not to do. Let's try that again, this time being mindful of which variables we actually want to include in our model.

```
[3]: df = pd.read_csv('titanic.csv')
df.head()
```

```
[3]: PassengerId  Survived  Pclass  \
0             1         0       3
1             2         1       1
2             3         1       3
3             4         1       1
4             5         0       3
```

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

```
[4]: x_feats = ['Pclass', 'Sex', 'Age', 'SibSp', 'Fare', 'Cabin', 'Embarked']
X = pd.get_dummies(df[x_feats], drop_first=True)
y = df['Survived']
X.head() # Preview our data to make sure it looks reasonable
```

```
[4]:
```

	Pclass	Age	SibSp	Fare	Sex_male	Cabin_A14	Cabin_A16	Cabin_A19	\
0	3	22.0	1	7.2500	1	0	0	0	
1	1	38.0	1	71.2833	0	0	0	0	
2	3	26.0	0	7.9250	0	0	0	0	
3	1	35.0	1	53.1000	0	0	0	0	
4	3	35.0	0	8.0500	1	0	0	0	

  

	Cabin_A20	Cabin_A23	...	Cabin_F G63	Cabin_F G73	Cabin_F2	Cabin_F33	\
0	0	0	...	0	0	0	0	
1	0	0	...	0	0	0	0	
2	0	0	...	0	0	0	0	
3	0	0	...	0	0	0	0	
4	0	0	...	0	0	0	0	

  

	Cabin_F38	Cabin_F4	Cabin_G6	Cabin_T	Embarked_Q	Embarked_S
0	0	0	0	0	0	1
1	0	0	0	0	0	0
2	0	0	0	0	0	1
3	0	0	0	0	0	1
4	0	0	0	0	0	1

[5 rows x 153 columns]

## 1.5 Normalization

Another important model tuning practice is to normalize your data. That is, if the features are on different scales, some features may impact the model more heavily than others. To level the playing field, we often normalize all features to a consistent scale of 0 to 1.

```
[5]: # Fill missing values
X = X.fillna(value=0)
for col in X.columns:
    # Subtract the minimum and divide by the range forcing a scale of 0 to 1
    ↪for each feature
    X[col] = (X[col] - min(X[col])) / (max(X[col]) - min(X[col]))

X.head()
```

```
[5]:
```

	Pclass	Age	SibSp	Fare	Sex_male	Cabin_A14	Cabin_A16	Cabin_A19	\
0	1.0	0.2750	0.125	0.014151	1.0	0.0	0.0	0.0	
1	0.0	0.4750	0.125	0.139136	0.0	0.0	0.0	0.0	
2	1.0	0.3250	0.000	0.015469	0.0	0.0	0.0	0.0	
3	0.0	0.4375	0.125	0.103644	0.0	0.0	0.0	0.0	
4	1.0	0.4375	0.000	0.015713	1.0	0.0	0.0	0.0	

  

	Cabin_A20	Cabin_A23	...	Cabin_F G63	Cabin_F G73	Cabin_F2	Cabin_F33	\
0	0.0	0.0	...	0.0	0.0	0.0	0.0	

1	0.0	0.0	...	0.0	0.0	0.0	0.0
2	0.0	0.0	...	0.0	0.0	0.0	0.0
3	0.0	0.0	...	0.0	0.0	0.0	0.0
4	0.0	0.0	...	0.0	0.0	0.0	0.0

	Cabin_F38	Cabin_F4	Cabin_G6	Cabin_T	Embarked_Q	Embarked_S
0	0.0	0.0	0.0	0.0	0.0	1.0
1	0.0	0.0	0.0	0.0	0.0	0.0
2	0.0	0.0	0.0	0.0	0.0	1.0
3	0.0	0.0	0.0	0.0	0.0	1.0
4	0.0	0.0	0.0	0.0	0.0	1.0

[5 rows x 153 columns]

## 1.6 Train-test split

```
[6]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
```

## 1.7 Fit a model

Fit an initial model to the training set. In scikit-learn, you do this by first creating an instance of the `LogisticRegression` class. From there, then use the `.fit()` method from your class instance to fit a model to the training data.

```
[9]: from sklearn.linear_model import LogisticRegression

logreg = LogisticRegression(fit_intercept=False, C=1e12, solver='liblinear')
model_log = logreg.fit(X_train, y_train)
model_log
```

```
[9]: LogisticRegression(C=1000000000000.0, fit_intercept=False, solver='liblinear')
```

## 1.8 Predict

Now that we have a model, let's take a look at how it performs.

```
[10]: y_hat_test = logreg.predict(X_test)
y_hat_train = logreg.predict(X_train)
```

```
[11]: import numpy as np
# We could subtract the two columns. If values are equal, difference will be
# zero. Then count number of zeros
residuals = np.abs(y_train - y_hat_train)
print(pd.Series(residuals).value_counts())
print(pd.Series(residuals).value_counts(normalize=True))
```

```
0    563
1    105
Name: Survived, dtype: int64
0    0.842814
1    0.157186
Name: Survived, dtype: float64
```

Not bad; our classifier was about 85% correct on our training data!

```
[12]: residuals = np.abs(y_test - y_hat_test)
      print(pd.Series(residuals).value_counts())
      print(pd.Series(residuals).value_counts(normalize=True))
```

```
0    174
1     49
Name: Survived, dtype: int64
0    0.780269
1    0.219731
Name: Survived, dtype: float64
```

And still about 80% accurate on our test data!

## 1.9 Summary

In this lesson, you took a more complete look at a data science pipeline for logistic regression, splitting the data into training and test sets and using the model to make predictions. You'll practice this on your own in the upcoming lab before having a more detailed discussion of more nuanced methods for evaluating a classifier's performance.